

Linux Privilege Escalation

Troy Defty

Lukasz Gogolkiewicz

who

- Lukasz


who

- Troy (@5ud0ch0p)
- 8 years in industry
 - 5.5 in UK, 2.5 AU
 - 7.5 red, 0.5 blue
- DEFCON 27, AusCERT, AISA, etc.

linux privesc

- privilege model
- recon
- auth weaknesses
- weak file permissions
- built-in escalation mechanisms && misconfiguration
- service misconfiguration
- artefact exploitation
- escaping restrictions
- * advanced (SELinux, LD_PRELOAD)

struct

- technique
 - hands-on
 - hints
 - review
- 
- 30 mins

- hands-on; 3 levels
 - intro
 - intermediate
 - annoying*

caveat

- some might seem simple
- some might seem strange
- some might seem confusing
- some might seem impossible

} welcome to
hacking!

- almost all challenges are representative of real scenarios we've encountered
- some (not many) are contrived to provide a means of practice

basics &&
reconnaissance

uid(0)

- 'root'
- highest user privilege on a *nix device
- as an attacker, high value:
 - read+write access to all data
 - credentials
 - database contents
 - defacement of webapps, etc
 - access to all functionality
 - repurposing device (!)
 - network pivot

uid(0)

- typical high privs:
 - device power control
 - control over peripherals and components
 - creating/starting/stopping services
 - user management
 - (un)installation of packages
 - device configuration
 - binding to privileged ports (1-1024)

usr vs krnl

ffffffff

kernel

- OS functionality:
 - interacting with hardware
 - handling interrupts
 - managing processes

user

- user applications:
 - text editing
 - web browsing
 - games

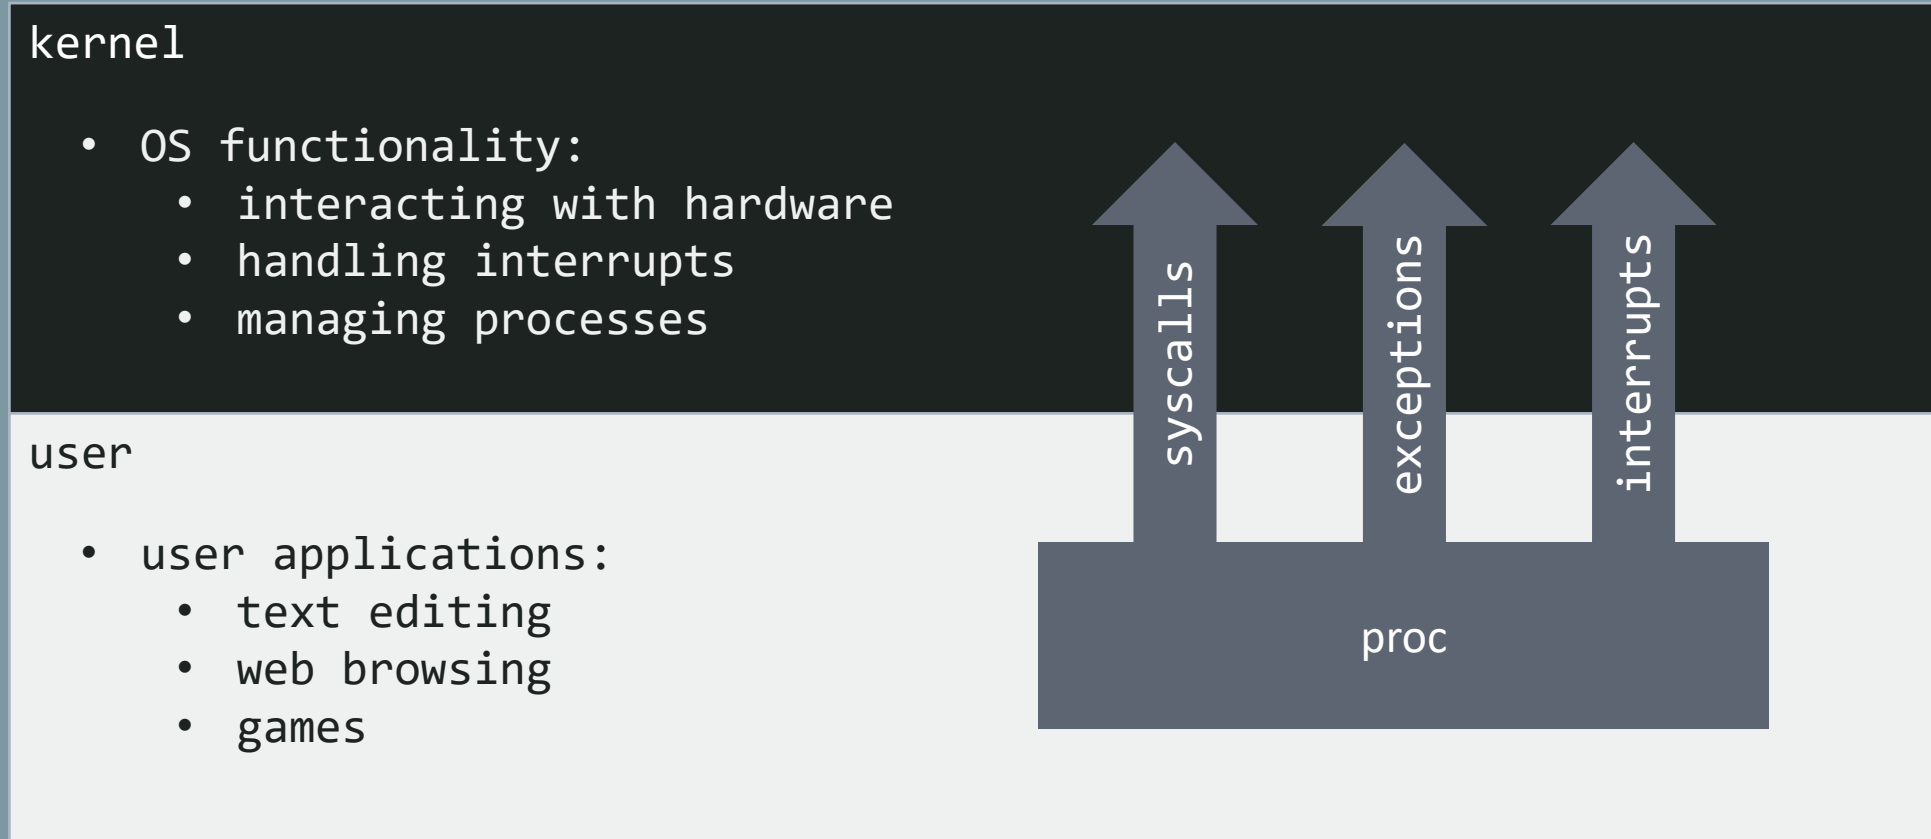
syscalls

exceptions

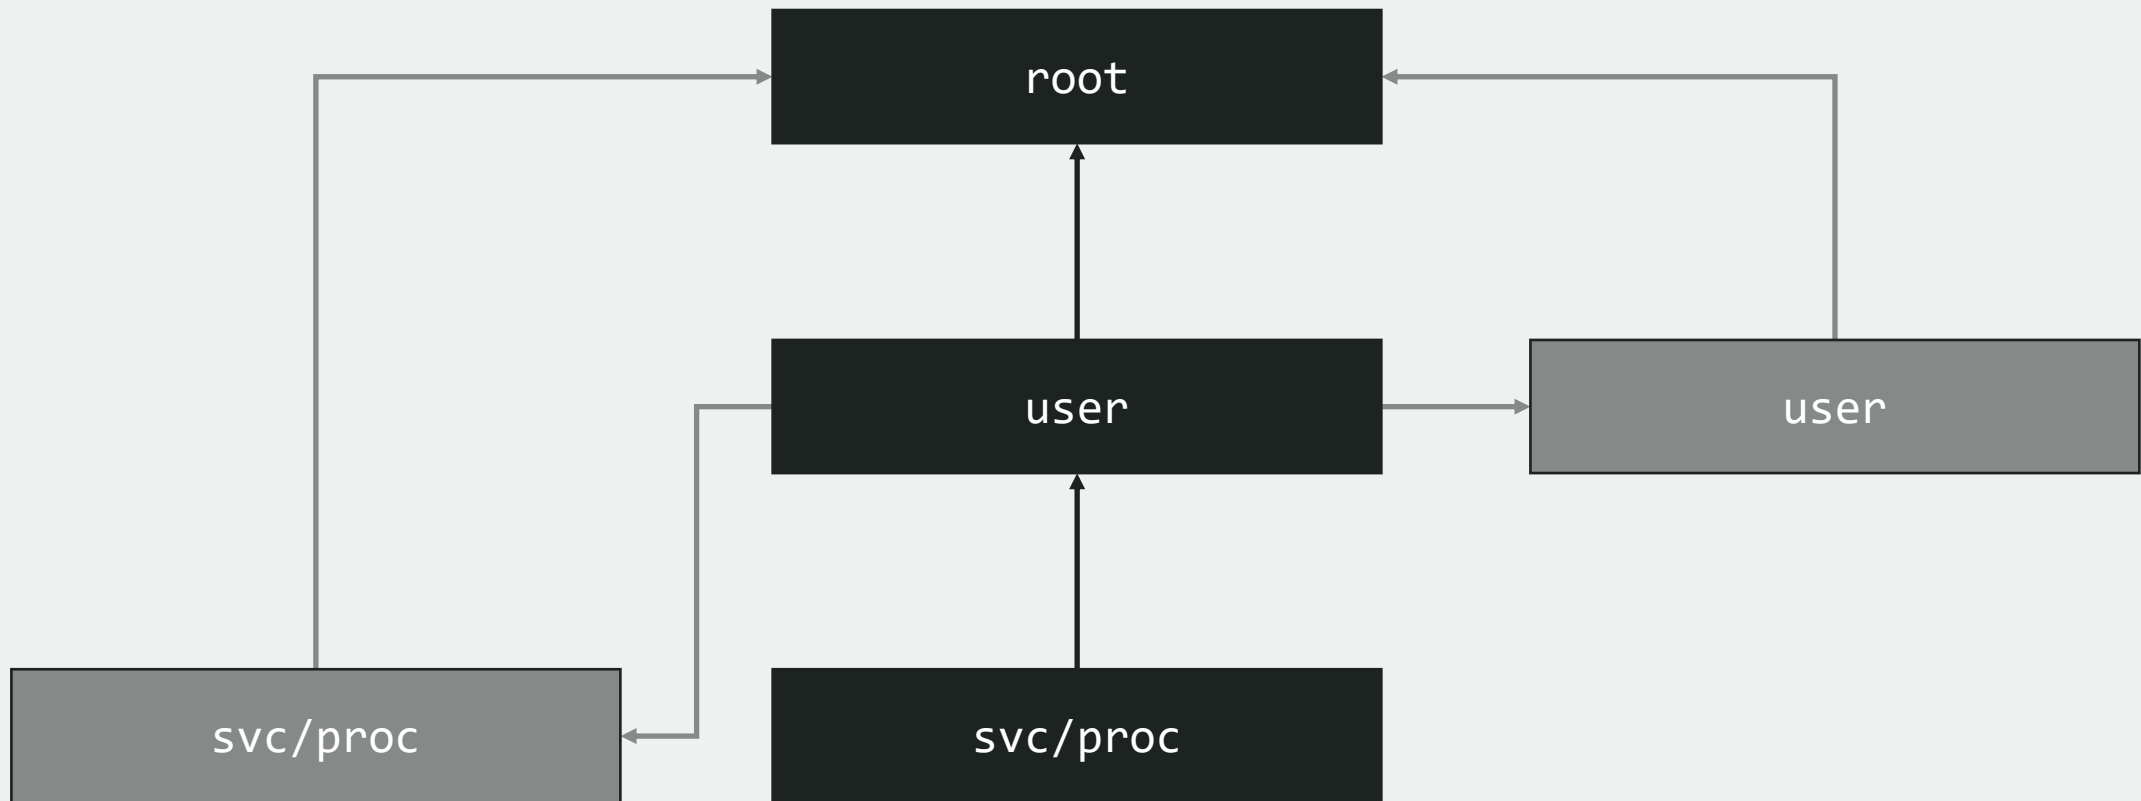
interrupts

proc

00000000



privilege escalation



users and authentication

- /etc/passwd, /etc/shadow, /etc/group
- pam
- svc specific mechanisms (.ssh/authorized_keys)
- su/sudo*

privs && access control

- “EvErYtHiNg Is A fIlE”
- discretionary access control - rwx
- mandatory access control – SELinux/apparmor, etc

perms & access control

- DAC : `rwXr-xr-x`

user	group	other
 <code>rwX</code>	 <code>r-x</code>	 <code>r-x</code>

```
-rw-r--r-- 1 root root 1734 Jun 14 08:58 /etc/passwd
```

```
drwx----- 2 dhcpcd dhcpcd 4096 Jun 14 08:58 dhcpcd
```

privs && access control

- MAC
- SELinux, apparmor, etc.
- rwx may not == rwx
- generally two modes:
 - report-only
 - enforce
- implementation specific:
 - SELinux* – contexts and policies
 - apparmor – path-dependent, mixing of modes

perms && access control

- `suid` : `rwsr-xr-x`
 - executes as user who owns file

```
-rwsr-xr-x 1 root root 63640 Feb 4 23:31 /usr/bin/passwd
```

- `sgid` : `rwxr-sr-x`
 - executes as group who owns file (*effective* gid)

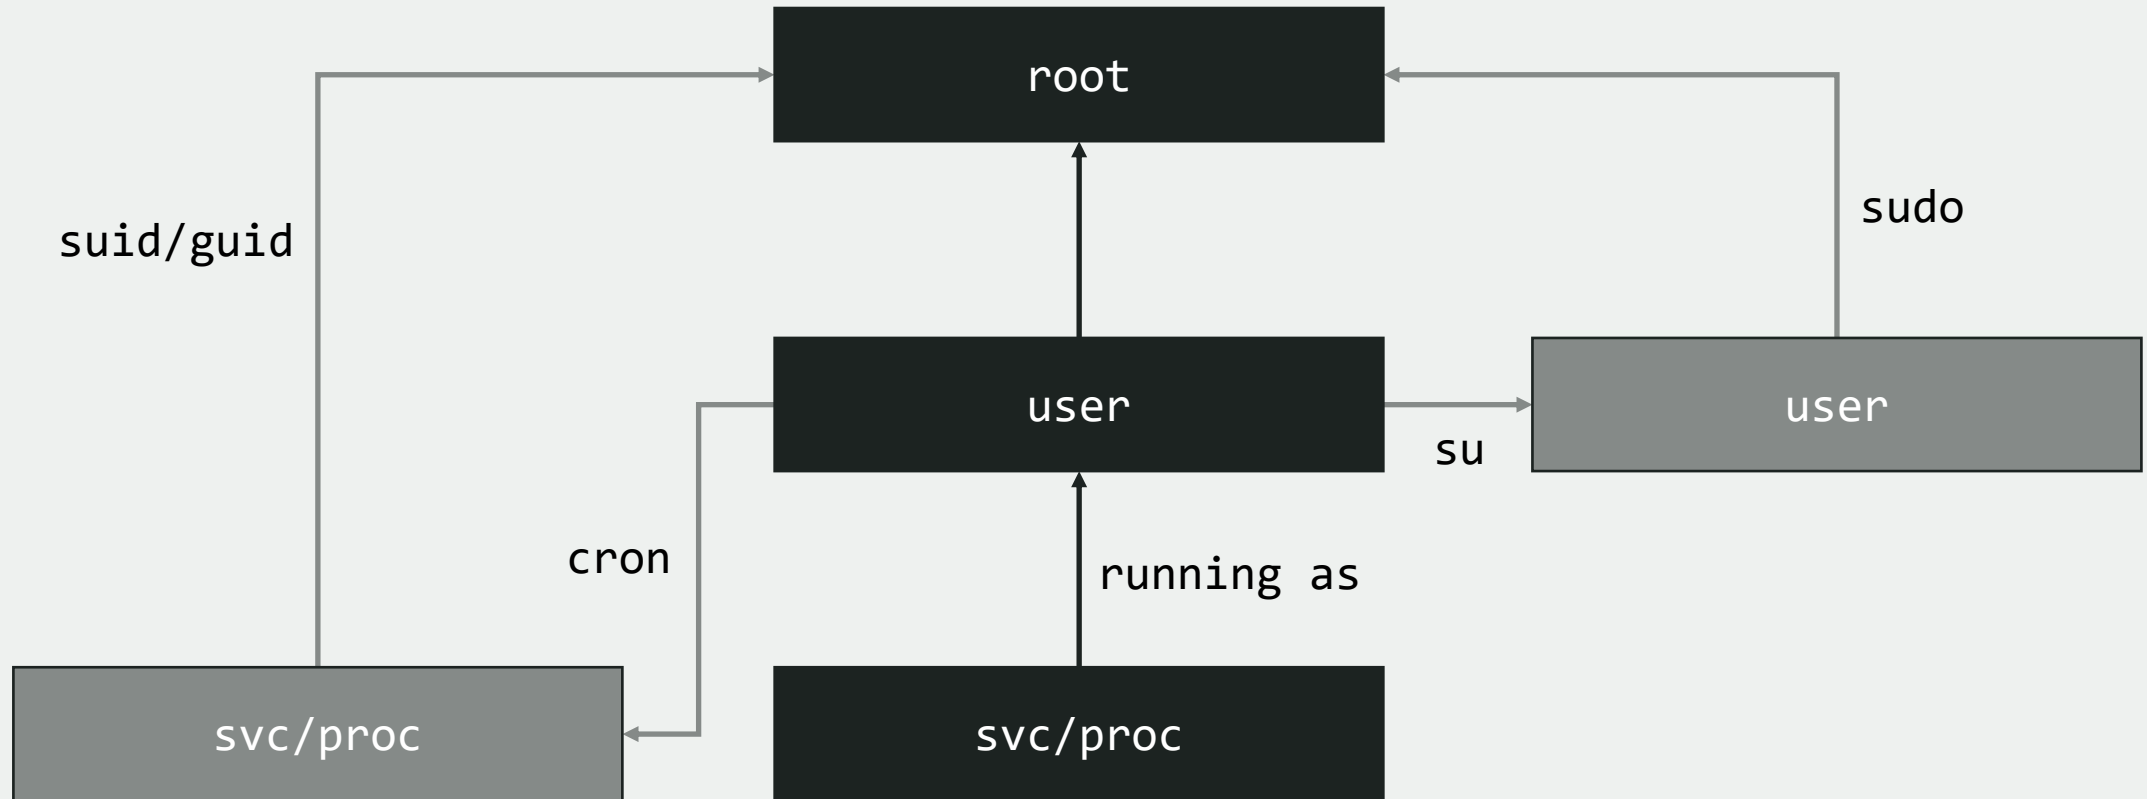
```
-rwxr-sr-x 1 root tty 34784 May 24 18:09 /usr/bin/wall
```


privs && svcs && procs

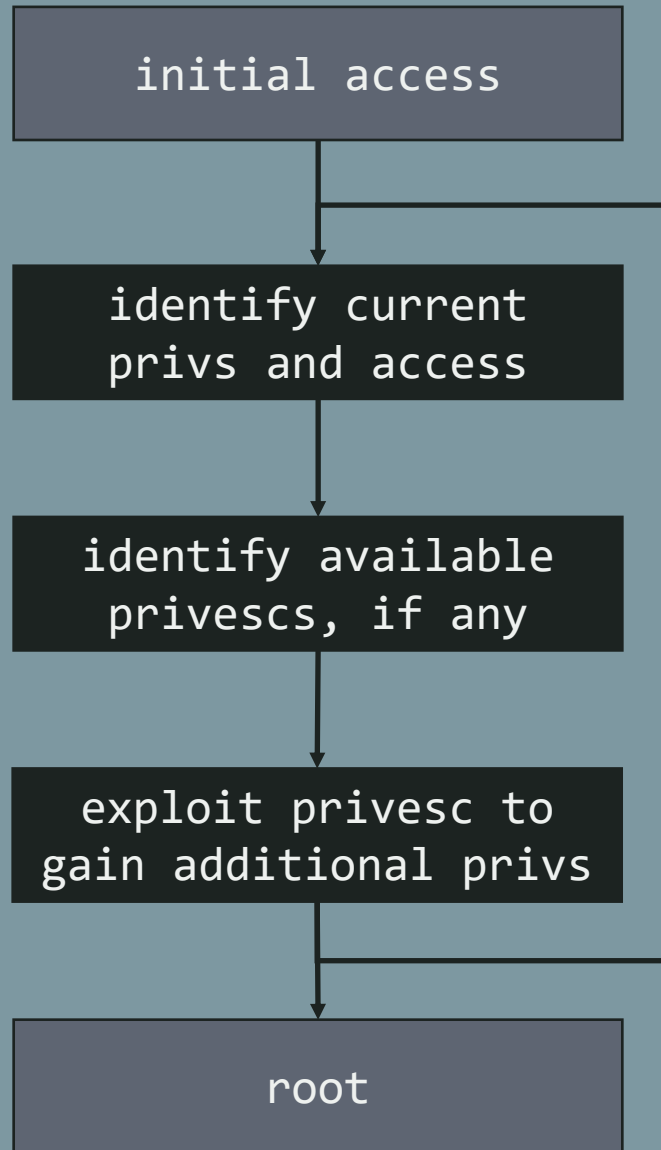
- everything runs as a 'user'
- svcs and procs need privs
 - running as priv. account
 - run as root
 - e.g. cron, systemd, etc
- svcs might also shed privs
 - run as low priv as possible
 - apache, nfs, etc.
 - wwwdata

} suid/sgid, e.g.

privilege model



method



recon

- users, groups
 - /etc/passwd, /etc/shadow, /etc/groups, etc.
- processes (&& configs)
 - ps/ss, documentation!
- services (&& configs)
 - netstat, systemctl, documentation!
- files && contents
 - ls, find, grep, strings, cat, etc.
- binaries && execution perms
 - ls, find, etc.
- access to all of the above

git && docker
good

quick git

```
git config --global core.autocrlf input
```

```
git clone https://github.com/5ud0ch0p/auscert2020-  
linux-privesc.git
```

docker setup && test

- docker setup
- lowpriv:lowpriv
- run some recon
 - get used to CentOS!

practical 0x00

- don't need automated tooling, but feel free
 - would strongly recommend to do without
- recommend:
 - ssh, cat, ls, find, grep, ps, netstat, etc.
 - documentation!
- don't diff the image!
 - yeah you'll solve the challenge
 - but we're here to learn!
 - likely can't ask the admin to roll the box back and then configure everything again in the real world!

authentication
weaknesses

authentication

- ssh, su/sudo, telnet, etc.
 - authentication often required
 - passwd
 - keys
 - pam
- auth data often mishandled, stored incorrectly, generally weak
 - “password”, “<username>”, “”
 - weak key lengths, permissions
 - plaintext storage
 - hash functions

authentication

- /etc/passwd

```
user1:x:1000:1000:User 1:/home/user1/:/bin/bash
```

```
user2:<hash>:1001:1001:User 2:/home/user2/:/bin/bash
```

- /etc/shadow

```
user1:$hashid$salt$hash:12345:0:90:10:::
```

```
user2:$hashid$salt$hash:11223:0:90:10:::
```

- /etc/group

```
group1:x:40:user1,user2
```

```
group2:x:41:user2
```

practical 0x01

- ssh as lowpriv
- ~/configure-privescs
 - to configure practical and difficulty
- get 'root'

review

- weak pwds
 - pwd-based auth only as good as pwd
- key OS files
- pwd storage mechanisms, hashing
- weak pwd storage

file permissions

file usage

- general user
 - scripts
 - backups
 - debug data
- service
 - configuration
 - authentication
- type
 - general plaintext
 - specific formats (zip, pcap, configs, etc.)
 - binary (programs, dump, etc.)
 - special (directories, socket, link, etc.)

file permissions

- file *and* directory
 - not all behave as expected!
 - s on directory?
 - suid *generally* ignored
 - sgid – new files and subdirs inherit groupID of dir, rather than usr
- rwx vs octal
 - think of binary bits

user	group	other
rwx	r-x	r-x
111	101	101
7	5	5
755		

dir structure

- /
 - /bin/ - 'core' binaries (ls, cat, cd, etc.)
 - /etc/ - configuration files
 - /home/ - user home directories
 - /tmp/ - temporary files
 - /usr/ - user-land programs and data
 - /var/ - things likely to change; logs, e.g.
 - (non-exhaustive)
- depends on distro!
- some have unique properties
 - /tmp/ - non-boot-persistent, 777

practical 0x02

- ssh as lowpriv
- get 'root'

review

- importance of:
 - acls
 - specific files/dirs.
- identifying:
 - files of interest
 - and ways of finding them

built-in escalation
mechanisms

perms plz

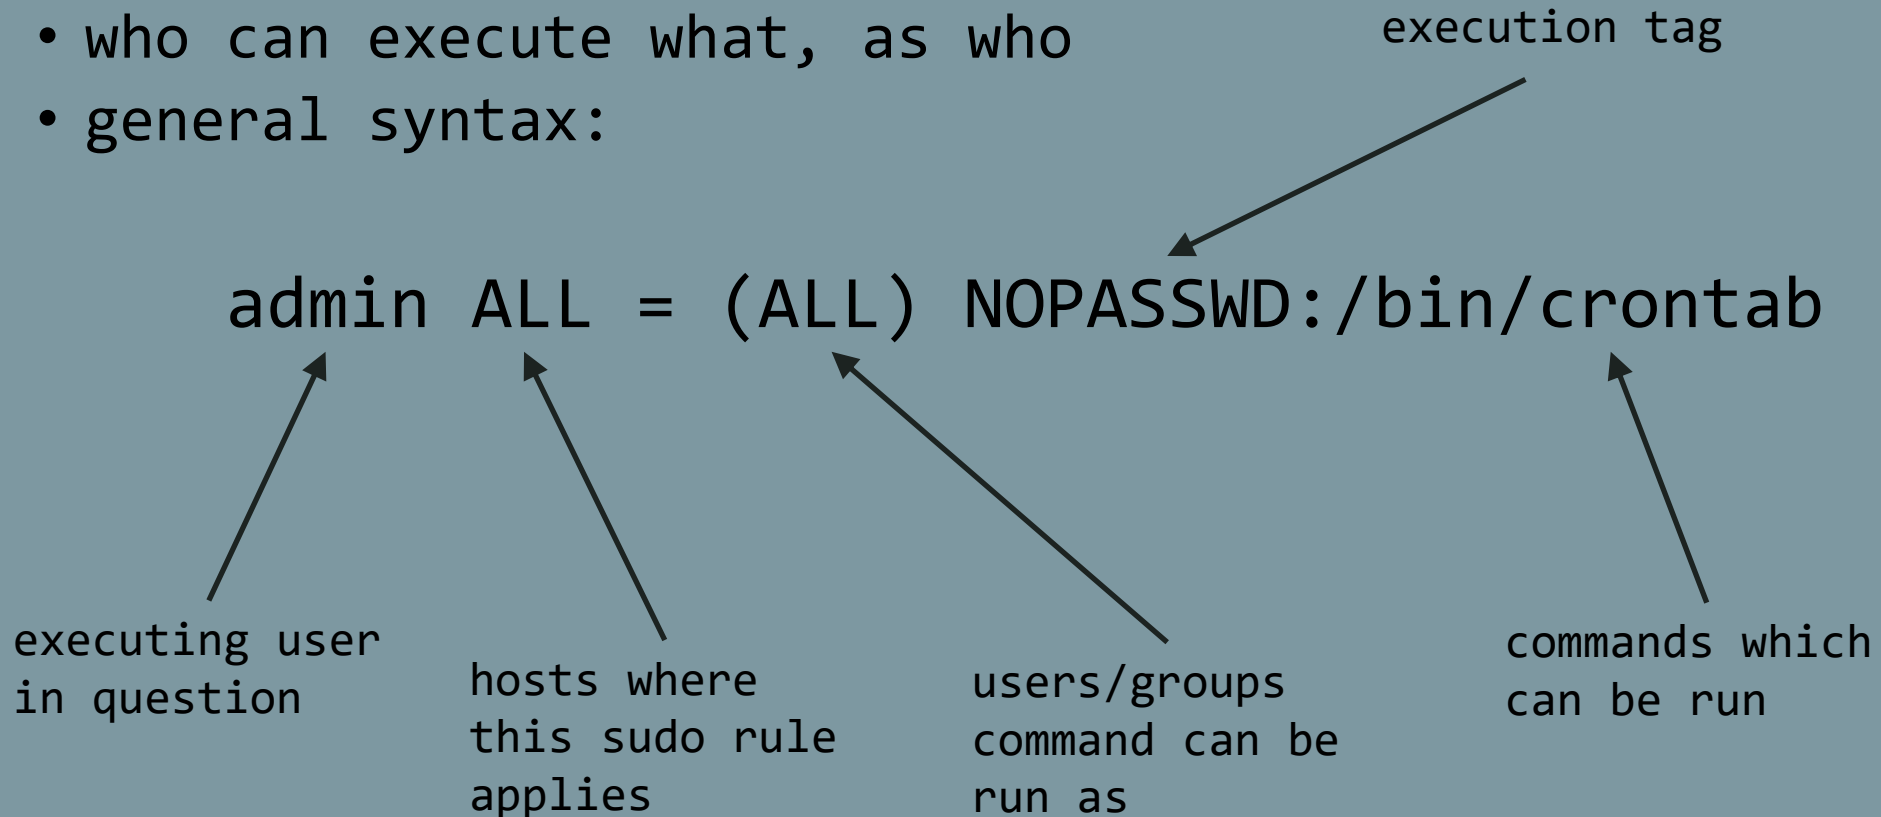
- usr needs additional permissions temporarily
 - doesn't need full root all the time
- usr:
 - (re)starting/stopping services
 - (un)installing packages
 - reboot machine

su(do)

- su
 - substitutes user and group IDs
 - spawns a shell (by default)
 - requires **target user** authentication success
- sudo
 - launches single command (by default)
 - requires **current user** authentication success
- 'su -c' ≈ 'sudo'
- 'su' ≈ 'sudo /bin/bash'

sudo

- /etc/sudoers
 - who can execute what, as who
 - general syntax:



sudo

- /etc/sudoers

- e.g.:

```
%wheel ALL = (ALL) ALL
```

```
%admins ALL = (ALL) NOPASSWD:/tmp/admin_tools/*
```

```
lowpriv localhost = /home/*/*/config
```


sudo

- `sudo -l`
- `/etc/sudoers.d/*`
- `man sudoers`

practical 0x03

- ssh as lowpriv
- get 'root'
- remember; file access permissions!

suid/sgid

- provides limited functionality using effective uID/gIDs
- usr:
 - change their password
 - changes /etc/shadow, owned by root:root
 - passwd needs to run as root!
 - mount a drive
 - privileged action
 - needs to run as root!
 - configure cron
 - /etc/cron* owned by root (as it contains all user crons!)
 - needs to run as root!

suid/sgid

- suid : rwsr-xr-x
 - executes as user who owns file

```
-rwsr-xr-x 1 root root 63640 Feb 4 23:31 /usr/bin/passwd
```

- sgid : rwxr-sr-x
 - executes as group who owns file

```
-rwxr-sr-x 1 root tty 34784 May 24 18:09 /usr/bin/wall
```

suid/sgid

- `find -perm:`
 - `-4000 = suid`
 - `-2000 = sgid`
 - `-6000 = both`
- don't need `sudo/su`

practical 0x04

- ssh as lowpriv
- get 'root'

review

- built-in escalation mechanisms
- dangers of sudo configs
- awareness of suid/sgid

artefacts and remnants

footprints in the sand

- usr actions leave traces
 - /tmp/, /home/usr/, etc.
 - histfile
 - syslog
- sysadmin processes can be insecure/incomplete in cleanup
 - user creation/deletion
 - software installation/removal

footprints in the sand

- recon is important!
- specific files often of interest
 - .bash_history
 - /var/log/*
- sysadmin processes leave remnants
 - *.bak
 - ./sysadmin.sh
 - mysql -u root -p <whatever>
 - orphaned userIDs
 - find / [-nouser|-nogroup]

practical 0x05

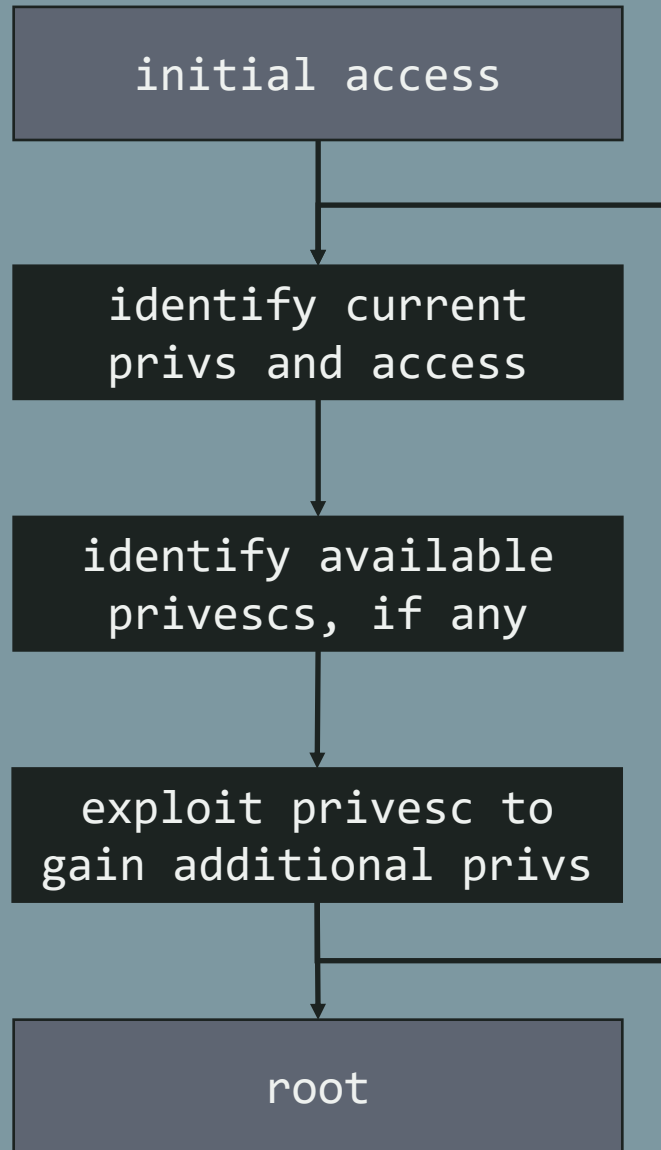
- ssh as lowpriv
- get 'root'

review

- sysadmin processes can introduce unintended consequences
- context dependent, but can be of use!

EOF

review



review

- privilege model
- recon
- auth weaknesses
- weak file permissions
- built-in escalation mechanisms && misconfiguration
- service misconfiguration
- artefact exploitation
- escaping restrictions
- * advanced (SELinux, LD_PRELOAD)

review

- importance of recon
- importance of methodology

?

[^]D

Linux Privilege Escalation

Troy Defty

Lukasz Gogolkiewicz

```
$DAY=$( (DAY+1) )
```

review

- privilege model
- recon
- auth weaknesses
- weak file permissions
- built-in escalation mechanisms && misconfiguration
- service misconfiguration
- artefact exploitation
- escaping restrictions
- * advanced (SELinux, LD_PRELOAD)

escaping restricted
execution environments

lockdown

- usr account might have limited permissions
- or a limited allowed command set
 - login places usr in limited execution environment
 - exploitation of service might grant access to limited features
- we want more privs/access!
 - escaping locked down environments
 - bypassing restrictions

breakout

- what are we running in?
- what can we do in our restricted environment?
- are there any documented escapes for this environment?
- is there a way we can fundamentally bypass the restrictions in place?
- what kind of input might break the restrictions?
 - `${}`, `${{}}`, `||`, `&&`, `<`, `>`, etc.

practical 0x06

- ssh as lowpriv
- get interactive command-line access (/bin/bash or similar) as root

review

- awareness of 'extra' functionality
- not always secure by default!
 - and this shouldn't be assumed!

service
misconfiguration

SVCS

- proc running (often backgrounded)
- run as a specific user
 - sometimes to shed perms (www-data)
 - sometimes because they need perms (root)

SVCS

- often provides a... service
 - db (mysql, postgresql, etc.)
 - http (apache, nginx, etc.)
 - scheduled jobs (cron, systemd, etc.)
 - remote management (telnet, SSH, etc.)
 - file sharing (ftp, etc.)
 - networking (dns, dhcp, etc.)
- configs can be complex and tricky
 - can introduce vulns, privescs!
 - mostly file-based (often /etc/)

identifying the privesc

- remember recon!
- what is running?
- what is it running as?
- where/how is it configured?
- documentation/manpages
- often distro-dependent

cron

- crontab
 - -l = list user crontab
 - -e = edit user crontab
- /etc/cron*:
 - crontab (the file)
 - cron.hourly/, cron.daily/, cron.weekly/, cron.hourly/
 - cron.d/ - contains system cronjobs for various users
 - cron.deny or cron.allow - crontab access controls

cron

- `/var/spool/cron/`:
 - file per user ('crontab')
 - editable via `crontab -e`
 - default perms `600`
- on CentOS (and in our container):
 - `/etc/cron*` generally used by `anacron`
 - `/var/spool/cron` generally used by `cron`

cron

- file syntax:

<m> <h> <day of month> <m> <day of week> <command>

*/20 * * * * zip -r logs.bak.zip /var/log/

2 5 * * * systemctl restart networking

0 9 9 6 * /root/start.sh

practical 0x07

- ssh as lowpriv

review

- services can be highly specific
- remember recon
- remember approach/methodology
- docs/manpages very helpful
- some behaviours not obvious

advanced:
shared objects

shared libs/objects

- compiled collections of functions, code, etc.
 - libc
 - libcrypt
 - libusb
- can be used by multiple programs
- given lib can have two “names”
 - library name (‘soname’) - libc.so.6
 - filename - /usr/lib/libc.so.6
- similar concept to DLLs in Windows

basic .so

- basic code structure of a shared library (in C):
 - header file (something.h)
 - source (something.c)

something.h

```
#ifndef ...  
#define ...  
  
extern void something(void);  
  
#endif
```

something.c

```
#include <stdio.h>  
  
void something(void) {  
    puts("I do something!");  
}
```

- then compile as a shared object

using our basic .so

- we can #include our SO similar to a core lib

libsomething.so

something.h

```
#ifndef ...  
#define ...  
  
extern void something(void);  
  
#endif
```

something.c

```
#include <stdio.h>  
  
void something(void) {  
    puts("I do something!");  
}
```

doathing.c

```
#include <stdio.h>  
#include "something.h"  
  
int main(void) {  
    puts("Lets do something");  
    something();  
    return 0;  
}
```

shared libs/objects

- SOs are **linked** during compilation, load time or run time
- list shared object dependencies for a given binary:
 - `ldd <binary>`
- list exported symbols from a lib:
 - `objdump -T /path/to/lib.so`
 - `nm -D /path/to/lib.so`
 - T prefix indicates export

linking

- **static**
 - all libs copied into main binary
 - all code, libs, etc. placed into memory at once by OS
 - once linked, libs are static and changes require recompilation
- **dynamic**
 - names of libs placed into binary
 - OS then loads main binary and libs separately at runtime
 - libs can change (within reason!) and main binary does not require recompilation

load order

- OS looks for dynamically-linked libs in various locations:
 - DT_RPATH in dynamic section of binary
 - LD_LIBRARY_PATH
 - DT_RUNPATH
 - /etc/ld.so.cache
 - /lib*
 - /usr/lib*

* Can also be /lib64, /usr/lib64

SO search path manipulation

- LD_LIBRARY_PATH
 - RPATH
 - LD_PRELOAD
-
- can be hacky solutions to dependency hell
 - often used for debugging
 - can be left behind after debugging!

LD_LIBRARY_PATH

- envvar
- *nix-specific (not all *nix, only some)
- contains colon-delimited list of dirs.
 - searched **before** typical search order directories
- how could this be problematic?

```
LD_LIBRARY_PATH=
```

```
DT_RPATH  
LD_LIBRARY_PATH  
DT_RUNPATH  
/etc/ld.so.cache  
/lib*  
/usr/lib*
```

```
LD_LIBRARY_PATH=/tmp/
```

```
/tmp/  
DT_RPATH  
LD_LIBRARY_PATH  
DT_RUNPATH  
/etc/ld.so.cache  
/lib*  
/usr/lib*
```

RPATH

- similar to LD_LIBRARY_PATH
- but compiled within binary
 - not dependent upon user envvars

-rpath=/path/to/something

- LD_RUN_PATH is envvar equivalent

RPATH

- `objdump -x /path/to/binary | grep RPATH`
- can we write to this location?
- can another user write to this location?
- `DT_RUNPATH` – similar `RPATH`, compiled into binary

LD_PRELOAD

- envvar
- lists SOs that override all normal shared objects
- often problematic with sudo:

Defaults

envkeep += LD_PRELOAD

- why?

practical 0x08

- ssh as lowpriv
- metasploit/msfvenom may be of help here
 - not necessary; can be done solely in C!
 - (if you have more time, would recommend giving it a go in C!)
 - why?
- automated tooling can make identifying this kind of privesc easier (sometimes!)

review

- .so
- load order important!
- LD_LIBRARY_PATH, RPATH, LD_PRELOAD

advanced:
selinux

caveat

- following section is a (very) high level summary
- selinux is complex
- section is to provide the absolute basics, should you find yourselves encountering selinux
- no practicals

in the real world

- available on multiple distros
 - (good) support: RedHat/CentOS, Fedora, Gentoo, et al.
 - in repos for: Ubuntu, Debian, (apparmor normally used) et al.
 - no official support: Arch (only kernel modules supported), et al.
- in use by default on, for example:



<https://source.android.com/security/selinux>

mandatory access control

- SELinux, apparmor, etc.
- rwx may not == rwx
- generally two modes:
 - report-only
 - enforce
- implementation specific:
 - SELinux* – contexts and policies
 - apparmor – path-dependent, mixing of modes

selinux

- confines user programs, system services
- ideally, minimize required privileges for a given proc
- enforced by the kernel
- no inherent concept of 'root'
 - root is subject to SELinux criteria!
- selinux users != linux users
 - OS maps linux users to selinux users
 - often used to also map to roles


access control mechanisms

- type enforcement
 - all subjects and objects are allocated a type
- role-based access control (RBAC)
 - selinux users associated to 1[+] roles
- multi-level security (MLS)
 - uses 'security level' to enforce selinux policies
 - "Top_Secret", "Confidential", etc.
- multi-category security (MCS)
 - categorises objects to enforce selinux policies
 - "Log_Files", "Customer_Data", etc.

access control and contexts

- contexts:
 - username
 - role
 - domain (or type)
 - level
- (almost) everything is assigned a label
 - network ports, files, hardware, etc.
 - access between labelled objects controlled by policy files, but can be manually adjusted (!)

contexts



```
-rw-rw-r-- lowpriv lowpriv standard_u:access_r:user_home_t:s0 notes.txt
```


inheritance

- default, context inheritance allowed
 - files created within a directory of context `dir_t` are also created with `dir_t`
 - child processes spawned from proc with `exec_t` also have `exec_t`
- different from DAC!
 - `dir = rw-`, file created file follows user default umask
- how could this be bad?

policies

- grouping of rules of explicit permissions, e.g.:
 - read/execute
 - bind/connect to a port
- typical policy consists of:
 - mapping file (.te)
 - “file contexts” file (.fc) [optional]
 - interface file (.if) [optional]
- compiled into .pp binaries to be loaded into kernel space
- collectively define a domain transition
- default policies exist, but specific

enforcement

- policy controls access between a labelled process and labelled objects
- different enforcement modes (non-exhaustive):
 - disabled – no policy loaded
 - permissive – warnings printed on policy violation
 - enforcing – access denied, logged, on policy violation
 - targeted (default on CentOS):
 - confines specific system processes (httpd, named, dhcpd, mysqld)
 - all other system and user processes run in unconfined domain
 - designed to protect key processes without harming UX

policy

https://wiki.gentoo.org/wiki/SELinux/Tutorials/Creating_your_own_policy_module_file

module ID and
version info



requirements
for this
policy



permission
definition



```
policy_module(diraccess, 1.0)
```

```
gen_require(`  
    type user_t;  
    type var_log_t;  
`)
```

```
allow user_t var_log_t:dir {getattr search open read}
```

<https://selinuxproject.org/page/ObjectClassesPerms>

domain transitions

- three conditions:
 - policy allows transition from origin domain to target
 - origin domain has execute on file
 - file context is defined as target domain entry point

```
type_transition <process> <origin_context> : <target_context>
```

```
type_transition backup_t backup_exec_t : fileaccess_t
```

tooling

- -Z
- getenforce/sestatus – selinux status
- chcon – similar to chmod/chown but temporary
- semanage – core selinux management, e.g.:
 - user role association
 - change security context of a target
 - proc permission management
- seinfo – query policy components
- ssh <user>/<selinux_role>@hostname
- many, many more

privesc...?

- looking for:
 - selinux permissive(!), or disabled(!!)
 - overly-permissive policy entries
 - overly-permissive type definitions
 - users/processes with incorrect context
 - type_transition of interest
- think about:
 - what can we achieve with our current selinux context?
 - how can we find this out?
 - do we need to try and access additional permissions?
 - if we are only interested in data, httpd_* might be sufficient!
 - but we need to be running in this or child processes; no migrating!

privesc...!

- if successful:
 - Disable selinux enforcing mode!
 - writing to etc_t/shadow_t
 - loading kernel modules
 - etc.

misc.

- /etc/selinux/config
- logging:
 - /var/log/messages
 - /var/log/audit/audit.log
 - /var/lib/setroubleshoot/se_troubleshoot_database.xml
 - systemd

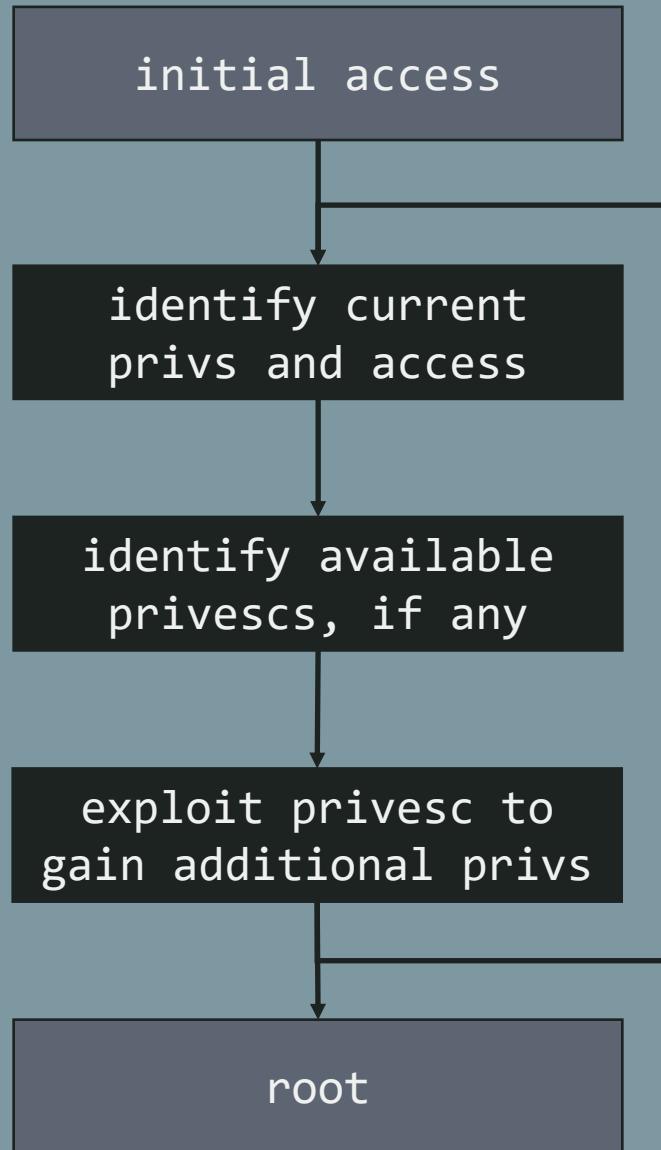
review

- selinux allocates a context to an object
- objects can only interact with their context
- can transition to other domains via policy

final challenges

EOF

review



review

- privilege model
- recon
- auth weaknesses
- weak file permissions
- built-in escalation mechanisms && misconfiguration
- service misconfiguration
- artefact exploitation
- escaping restrictions
- advanced (SELinux, LD_PRELOAD)

review

- importance of recon
- importance of methodology

?

[^]D